

# Learning ElasticSearch

**Anurag Patel — Red Hat**

Fifth Elephant 2013, Bangalore.

**<http://xinh.org/5el>**

Also available at

# ElasticWho?

ElasticSearch is a flexible and powerful open source, distributed real-time search and analytics engine.




# Features

- Real time analytics
- Distributed
- High availability
- Multi tenant architecture
- Full text
- Document oriented
- Schema free
- RESTful API
- Per-operation persistence

# Distributed

Start small and scale horizontally out of the box. For more capacity, just add more nodes and let the cluster reorganize itself.

# High Availability



ElasticSearch clusters detect and remove failed nodes, and reorganize themselves.

# Multi Tenancy

```
$ curl -XPUT http://localhost:9200/people  
$ curl -XPUT http://localhost:9200/gems  
$ curl -XPUT http://localhost:9200/gems/document/pry-0.5.9  
$ curl -XGET http://localhost:9200/gems/document/pry-0.5.9
```

A cluster can host multiple indices which can be queried independently, or as a group.

# Document Oriented

```
{
  "_id": "pry-0.5.9",
  "_index": "gems",
  "_source": {
    "authors": [
      "John Mair (banisterfiend)"
    ],
    "autorequire": null,
    "bindir": "bin",
    "cert_chain": [],
    "date": "Sun Feb 20 11:00:00 UTC 2011",
    "default_executable": null,
    "description": "attach an irb-like session to any object at runtime",
    "email": "jrmair@gmail.com"
  }
}
```

Store complex real world entities in Elasticsearch as structured JSON documents.



# RESTful API

Almost any operation can be performed using a simple RESTful interface using **JSON** over **HTTP**.

- curl -X GET
- curl -X PUT
- curl -X POST
- curl -X DELETE

# Apache Lucene

ElasticSearch is built on top of Apache Lucene. Lucene is a high performance, full-featured Information Retrieval library, written in Java.



# ElasticSearch Terminology

# Document

```
$ curl -XGET http://localhost:9200/gems/document/pry-0.5.9
```

```
{
  "_id": "pry-0.5.9",
  "_index": "gems",
  "_source": {
    "authors": [
      "John Mair (banisterfiend)"
    ],
    "autorequire": null,
    "bindir": "bin",
    "cert_chain": [],
    "date": "Sun Feb 20 11:00:00 UTC 2011",
    "default_executable": null,
    "description": "attach an irb-like session to any object at runtime",
    "email": "jrmair@gmail.com",
    "executables": [
      "pry"
    ],
    "extensions": [],
    "extra_rdoc_files": [],
    "files": [
      "lib/pry/commands.rb"
    ]
  }
}
```

In Elasticsearch, everything is stored as a Document. Document can be addressed and retrieved by querying their attributes.

# Document Types

Lets us specify document properties, so we can differentiate the objects.

# Shard


Each Shard is a separate *native* Lucene Index. Lets us overcome RAM limitations, hard disk capacity.

# Replica

An exact copy of primary Shard. Helps in setting up HA, increases query throughput.

# Index

ElasticSearch stores its data in logical Indices. Think of a table, collection or a database.



An Index has atleast 1 primary Shard, and 0 or more Replicas.



# Cluster

A collection of cooperating Elasticsearch nodes. Gives better availability and performance via Index Sharding and Replicas.

# ElasticSearch Workshop

# Download and start

Download Elasticsearch from  
<http://www.elasticsearch.org/download>

```
# service elasticsearch start
```

```
# /etc/init.d/elasticsearch start
```

```
# ./bin/elasticsearch -f
```

# ElasticSearch Plugins

A site plugin to view contents of ElasticSearch cluster.

```
# cd /usr/share/elasticsearch  
# ./bin/plugin -install mobz/elasticsearch-head
```

```
# cd /opt/elasticsearch-0.90.2  
# ./bin/plugin -install mobz/elasticsearch-head
```

Restart ElasticSearch. Plugins are detected and loaded on service startup.

# elasticsearch-head

The screenshot displays the ElasticSearch Head web interface in a browser window. The address bar shows `localhost:9200/_plugin/head/`. The main header includes the **ElasticSearch** logo, a connection URL `http://localhost:9200/`, a **Connect** button, and the **Enchantress** status **cluster health: green (1, 0)**. Below the header is a navigation bar with tabs: Overview, Browser, Structured Query [+], Any Request [+], Info, Status, Node Stats, Cluster Nodes, Cluster State, and Cluster Health. The **Cluster Overview** tab is active, showing a **New Index** button and a **Refresh** dropdown.

The main content area displays two indices: **aadhaar** and **gems**. Each index has a size and document count (both 0), and buttons for **Info** and **Actions**. Below the index information, a table shows the distribution of shards across nodes. The **Enchantress** node (ID: `SqF5EPSMQMypn8n1_7z31A`, IP: `inet[/192.168.131.187:9200]`) is highlighted in yellow and has 6 shards assigned (0-5). The **Unassigned** section shows 5 unassigned shards (0-4) for the **gems** index.

Index	Shard	Node
aadhaar	0	Enchantress
	1	Enchantress
	2	Enchantress
	3	Enchantress
	4	Enchantress
	5	Enchantress
gems	0	Unassigned
	1	Unassigned
	2	Unassigned
	3	Unassigned
	4	Unassigned

# RESTful interface

```
$ curl -XGET 'http://localhost:9200/'
```

```
{
  "ok" : true,
  "status" : 200,
  "name" : "Drake, Frank",
  "version" : {
    "number" : "0.90.2",
    "snapshot_build" : false,
    "lucene_version" : "4.3.1"
  },
  "tagline" : "You Know, for Search"
}
```

# Create Index

```
$ curl -XPUT 'http://localhost:9200/gems'
```

```
{  
  "ok":true,  
  "acknowledged":true  
}
```

# Cluster status

```
$ curl -XGET 'localhost:9200/_status'
```

```
{"ok":true,"_shards":{"total":20,"successful":10,"failed":0},  
"indices":{"gems":{"index":{"primary_size":"495b","primary_size_in_bytes":495,  
"size":"495b","size_in_bytes":495},"translog":{"operations":0},  
"docs":{"num_docs":0,"max_doc":0,"deleted_docs":0},"merges":  
{"current":0,"current_docs":0,"current_size":"0b","current_size_in_bytes":0,  
"total":0,"total_time":"0s","total_time_in_millis":0,"total_docs":0,  
"total_size":"0b","total_size_in_bytes":0},  
...  
...  
...
```



# Pretty Output

```
$ curl -XGET 'localhost:9200/_status?pretty'
```

```
$ curl -XGET 'localhost:9200/_status' | python -mjson.tool
```

```
$ curl -XGET 'localhost:9200/_status' | json_reformat
```

```
{
  "ok": true,
  "_shards": {
    "total": 20,
    "successful": 10,
    "failed": 0
  },
  "indices": {
    "gems": {
      "index": {
        "primary_size": "495b",
        "primary_size_in_bytes": 495,
        "size": "495b",
        "size_in_bytes": 495
      },
      ...
    }
  }
}
```

# Delete Index

```
$ curl -XDELETE 'http://localhost:9200/gems'
```

```
{  
  "ok":true,  
  "acknowledged":true  
}
```

# Create custom Index

```
{
  "settings" : {
    "index" : {
      "number_of_shards" : 6,
      "number_of_replicas" : 0
    }
  }
}
```

```
$ curl -XPUT 'http://localhost:9200/gems' -d @body.json
```

```
{
  "ok":true,
  "acknowledged":true
}
```

# Index a document

```
{
  "name": "pry",
  "platform": "ruby",
  "rubygems_version": "1.5.2",
  "description": "attach an irb-like session to any object at runtime",
  "email": "anurag@example.com",
  "has_rdoc": true,
  "homepage": "http://banisterfiend.wordpress.com"
}
```

```
$ curl -XPOST 'http://localhost:9200/gems/test/' -d @body.json
```

```
{
  "ok": true,
  "_index": "gems",
  "_type": "test",
  "_id": "lsJgxiwET6eg",
  "_version": 1
}
```

# Get document

```
$ curl -XGET 'http://localhost:9200/gems/test/lsJgxiwET6eg' | python -mjson.tool
```

```
{
  "_id": "lsJgxiwET6eg",
  "_index": "gems",
  "_source": {
    "description": "attach an irb-like session to any object at runtime",
    "email": "anurag@example.com",
    "has_rdoc": true,
    "homepage": "http://banisterfiend.wordpress.com",
    "name": "pry",
    "platform": "ruby",
    "rubygems_version": "1.5.2"
  },
  "_type": "test",
  "_version": 1,
  "exists": true
}
```

# Index another document

```
{
  "name": "grit",
  "platform": "jruby",
  "rubygems_version": "2.5.0",
  "description": "Ruby library for extracting information from a git repository.",
  "email": "mojombo@github.com",
  "has_rdoc": false,
  "homepage": "http://github.com/mojombo/grit"
}
```

```
$ curl -XPOST 'http://localhost:9200/gems/test/' -d @body.json
```

```
{
  "ok": true,
  "_index": "gems",
  "_type": "test",
  "_id": "ijU0Hi2cQc2",
  "_version": 1
}
```

# Custom Document IDs

```
{
  "name": "grit",
  "platform": "jruby",
  "rubygems_version": "2.5.1",
  "description": "Ruby library for extracting information from a git repository.",
  "email": "mojombo@github.com",
  "has_rdoc": false,
  "homepage": "http://github.com/mojombo/grit"
}
```

```
$ curl -XPUT 'http://localhost:9200/gems/test/grit-2.5.1' -d @body.json
```

```
{
  "ok": true,
  "_index": "gems",
  "_type": "test",
  "_id": "grit-2.5.1",
  "_version": 1
}
```

IDs are unique across Index. Composed of DocumentType and ID.

# Document Versions

```
$ curl -XPUT 'http://localhost:9200/gems/test/grit-2.5.1' -d @body.json
```

```
{  
  "ok":true,  
  "_index":"gems",  
  "_type":"test",  
  "_id":"grit-2.5.1",  
  "_version":2  
}
```



# Searching Documents

```
{
  "query": {
    "term": {"name": "pry"}
  }
}
```

```
$ curl -XPOST http://localhost:9200/gems/_search -d @body.json | python -mjson.tool
```

```
{
  "_shards": {
    "failed": 0,
    "successful": 6,
    "total": 6
  },
  "hits": {
    "hits": [
      {
        "_id": "MWkKgzsMRgK",
        "_index": "gems",
        "_score": 1.4054651,
        "_source": {
          "description": "attach an irb-like session to any object at runtime",
          "email": "anurag@example.com",
          "has_rdoc": true,
          "homepage": "http://banisterfiend.wordpress.com",
          "name": "pry",
          "platform": "ruby",
          "rubygems_version": "1.5.2"
        }
      }
    ]
  }
}
```

# Counting Documents

```
{  
  "term": {"name": "pry"}  
}
```

```
$ curl -XGET http://localhost:9200/gems/test/_count -d @body.json
```

```
{  
  "_shards": {  
    "failed": 0,  
    "successful": 6,  
    "total": 6  
  },  
  "count": 1  
}
```

# Update a Document

```
{  
  "doc": {  
    "platform": "macruby"  
  }  
}
```

```
$ curl -XPOST http://localhost:9200/gems/test/grit-2.5.1/_update -d @body.json
```

```
{  
  "ok":true,  
  "_index":"gems",  
  "_type":"test",  
  "_id":"grit-2.5.1",  
  "_version":4  
}
```

The partial document is merged using simple recursive merge.

# Update via Script

```
{  
  "script" : "ctx._source.platform = vm_name",  
  "params" : {  
    "vm_name" : "rubinius"  
  }  
}
```

```
$ curl -XPOST http://localhost:9200/gems/test/grit-2.5.1/_update -d @body.json
```

```
{  
  "ok":true,  
  "_index":"gems",  
  "_type":"test",  
  "_id":"grit-2.5.1",  
  "_version":5  
}
```

# Delete Document

```
$ curl -XDELETE 'http://localhost:9200/gems/test/grit-2.5.1'
```

```
{  
  "ok":true,  
  "found":true,  
  "_index":"gems",  
  "_type":"test",  
  "_id":"grit-2.5.1",  
  "_version":6  
}
```

# Put Mapping

```
{
  "gem" : {
    "properties" : {
      "name" :      {"type" : "string", "index": "not_analyzed"},
      "platform" :  {"type" : "string", "index": "not_analyzed"},
      "rubygems_version" : {"type" : "string", "index": "not_analyzed"},
      "description" : {"type" : "string", "store" : "yes"},
      "has_rdoc" :   {"type" : "boolean"}
    }
  }
}
```

```
$ curl -XPUT 'http://localhost:9200/gems/gem/_mapping' -d @body.json
```

```
$ curl -XGET 'http://localhost:9200/gems/_mapping' | python -mjson.tool
```

# Index Document with Mapping

```
{
  "name": "grit",
  "platform": "ruby",
  "rubygems_version": "2.5.1",
  "description": "Ruby library for extracting information from a git repository.",
  "email": "mojombo@github.com",
  "has_rdoc": false,
  "homepage": "http://github.com/mojombo/grit"
}
```

```
$ curl -XPUT 'http://localhost:9200/gems/gem/grit-2.5.1' -d @body.json
```

```
{
  "ok": true,
  "_index": "gems",
  "_type": "gem",
  "_id": "grit-2.5.1",
  "_version": 1
}
```

# Matching documents

```
{  
  "query": {  
    "match" : {  
      "description" : "git repository"  
    }  
  }  
}
```

```
$ curl -XPOST http://localhost:9200/gems/gem/_search -d @body.json
```



# Highlighting

```
{
  "query": {
    "match" : {
      "description" : "git repository"
    }
  },
  "highlight" : {
    "fields" : {
      "description" : {}
    }
  }
}
```

```
$ curl -XPOST http://localhost:9200/gems/gem/_search -d @body.json
```

```
"highlight": {
  "description": [
    "Ruby library for extracting information from a <em>git</em> <em>repository</em>."
  ]
}
```

# Search Facets

```
{
  "query": { "match_all" : {} },
  "facets" : {
    "gem_names" : {
      "terms" : { "field": "name" }
    }
  }
}
```

```
$ curl -XPOST http://localhost:9200/gems/_search -d @body.json
```

```
...
  "facets": {
    "gem_names": {
      "_type": "terms",
      "missing": 0,
      "other": 0,
      "terms": [
        {
          "count": 2,
          "term": "pry"
        },
        {
          "count": 2,
          "term": "grit"
        },
        {
          "count": 1,
          "term": "abc"
        }
      ]
    },
    "total": 5
  }
}
```

# **(Lab)**

# **Analyzing Aadhaar's Datasets**

# Download Public Dataset

Download from Aadhaar Public Data Portal at  
[\*\*https://data.uidai.gov.in\*\*](https://data.uidai.gov.in)

# Download Tools

```
$ git clone https://github.com/gnurag/aadhaar
```

# Prepare Data & Configure

```
# gem install yajl-ruby tire activesupport  
  
$ git clone https://github.com/gnurag/aadhaar  
$ cd aadhaar/data  
$ unzip UIDAI-ENR-DETAIL-20121001.zip  
$ cd ../bin  
$ vi aadhaar.rb
```

# Configuration

```
AADHAAR_DATA_DIR = "/path/to/aadhaar/data"  
ES_URL            = "http://localhost:9200"  
ES_INDEX          = 'aadhaar'  
ES_TYPE           = "UID"  
BATCH_SIZE        = 1000
```

# Index

\$ ruby aadhaar.rb



# Running Examples

```
$ curl -XPOST http://localhost:9200/aadhaar/UID/_search -d  
@template.json | python -mjson.tool
```

# **Additional Notes**

# Index Aliases

Group multiple Indexes, and query them together.

```
curl -XPOST 'http://localhost:9200/_aliases' -d '{
  "actions" : [
    { "add" : { "index" : "index1", "alias" : "master-alias" } }
    { "add" : { "index" : "index2", "alias" : "master-alias" } }
  ]
}'
```

```
curl -XPOST 'http://localhost:9200/_aliases' -d '{
  "actions" : [
    { "remove" : { "index" : "index2", "alias" : "master-alias" } }
  ]
}'
```

# Document Routing

Control which Shard the document will be placed and queried from.

# Parents & Children

```
$ curl -XPUT http://localhost:9200/gems/gem/roxml?parent=rexml -d '{  
  "tag" : "something"  
}'
```

# Custom Analyzers

# Boosting Search Results

# ElasticSearch Ecosystem

A wide range of site plugins, analyzers, river plugins available from the community.



# THE END

@gnurag / github